Shubham Joshi

shubham05373@gmail.com | 8983230919 | inferno-721.vercel.app

linkedin.com/in/shubham-joshi-721inferno | github.com/Inferno-721

Education

Indian Institute of Information Technology, Design and Manufacturing, Jabalpur, Btech — Dec 2021 – July 2025

- GPA: 8.0/10.0

Work Experience

AI Engineer, String Ventures—Michigan, USA — Jan 2025 – present

- Architected a hierarchical multi-agent LLM orchestration system on Google ADK coordinating 8+ specialized agents (intake router, forecast sub-orchestrator, hazard/lightning/hail specialists) with LiteLLM cross-provider failover and a 3-tier complexity router that maps each request to an optimal model + reasoning-effort tier, plus a deterministic fast-path gate that collapses simple queries to a single LLM call to cut latency and token cost.
- Designed the LLM function-calling / tool layer that grounds agent outputs in deterministic computation, having agents select transforms while a registry-driven engine computes 108 weather attributes across 7 summarizer types over live FLASH Weather AI forecast/lightning/hail data—eliminating numeric hallucination in generated responses.
- Led an eval-first prompt-engineering workflow, building a 60-query natural-language eval harness with deterministic pass/fail checks and tuning agent prompts, routing signals, and complexity tiers to a 60/60 pass rate, alongside LLM-based conversation memory that summarizes multi-turn context for cross-session continuity.
- **Tech Stack:** Python, FastAPI, Google ADK, LiteLLM, Prompt Engineering, Function Calling, Eval Harnesses.

Data Science Intern (Part-time), InfoOrigin—Noida, India — Dec 2024–July 2025

- Built and productionized an AI-powered NL2SQL/RAG platform for natural language querying and report generation across 200+ database schemas, serving 70+ queries daily, using Qdrant, semantic search, and LLMs for query routing.
- Automated end-to-end analytics workflows with streaming APIs, multi-session conversation management, and intelligent context retrieval, reducing manual data analysis effort by over **80%** while improving reporting reliability.
- **Tech Stack:** Python, FastAPI, SQL, RAG, Qdrant, Vector Databases, LLMs, Prompt Engineering, Embeddings, Streaming APIs, Automation, Session Management

Projects

CRM MCP Server

- Built a production MCP (Model Context Protocol) server in Go that lets external AI agents connect to a multi-tenant CRM — secured with OAuth 2.0 and role-based tool permissions — enabling agents to run a full 5-step prospecting workflow end to end.
- Added human-in-the-loop safeguards: agent research is stored outside the CRM first and only written in after a person reviews and approves it, with automated reports flagging incomplete records — cutting manual data entry while keeping every CRM write human-approved.
- **Tech Stack:** Go, Model Context Protocol (MCP), OAuth 2.0, REST APIs, Agentic Workflows

Skills

Languages: Python, SQL

Tools: Jupyter Notebook, VS Code, PyCharm, Git, GitHub, Metabase, Power BI

Frameworks: LlamaIndex, LangChain, LangGraph, FastAPI, HuggingFace, Streamlit

Machine Learning & AI: Machine Learning, Deep Learning, Generative AI, Transformers, LLM Agents

Achievements

- **JENESYS Scholar**—Represented India in Japan as 1 of 30 students selected nationwide for academic and leadership excellence.
- **AIR 1 – Data Analysis (DDC 2023):** Secured **1st place** among 50+ national teams.